





UNDERSTANDING CONTAINER IMAGES

DATA CENTRE - DEVOPS CHAPTER

ALVARO REVUELTA



PRESENTATION STRUCTURE

1. Introduction
2. Image Structure
3. How layers are generated
4. Optimize images
5. Conclusion



1. INTRODUCTION

- The presentation will be intercalated with a demo using the tool Dive and docker command line.

Motivation

- Bridging Knowledge: **Complex technical stack.**
- Understanding the importance of formation in virtualization and container technology for **everyone.**




1. INTRODUCTION

- Images are **immutable**. Once an image is created, it can't be modified.
- Container images are composed of **layers**. Each layer represents a set of file system changes that add, remove, or modify files.



2. IMAGE STRUCTURE

- Images consist of **multiple layers** stacked on top of each other.
- Layers are **file system changes** (e.g., file additions, deletions, updates).
- Layers are **read-only**; only the top container layer is **writable** at runtime.
- Layers are merged to form a **Unified File System** (Full explanation would need another presentation).



(Source: UnionFS : A File System of a Container)



3. HOW IMAGES ARE GENERATED

- Layers are **reusable** in subsequent builds, because they are cached.
- **FROM:** Setting the base image (starting point)
- **RUN:** Run a command (e.g., installing software)
- **COPY/ADD:** Adding files
- **CMD/ENTRYPOINT:** Setting the default command for running the container
- RUN, ADD and COPY generate new layers.

Image hierarchy			
FROM	debian:12, 12.7, bookworm, bookworm-20240926, latest	49.56 MB	1
FROM	buildpack-deps:bookworm-curl, curl, stable-curl	0 B	1
FROM	buildpack-deps:bookworm-scm, scm, stable-scm	24.05 MB	1
FROM	buildpack-deps:bookworm	64.39 MB	1
ALL	node:latest	211.27 MB	1
Layers (13)			
0	ADD file:087f68d5558e06c7160c9322582925635e7539a7702413828357c28c77f6f345 in /	49.56 MB	1
1	CMD ["bash"]	0 B	1
2	/bin/sh -c set -eux; apt-get update; apt-get install -y --no-install-recommends ca-certificates curl ...	24.05 MB	1
3	/bin/sh -c set -eux; apt-get update; apt-get install -y --no-install-recommends git mercurial opens...	64.39 MB	1
4	/bin/sh -c set -eux; apt-get update; apt-get install -y --no-install-recommends autoconf automake b...	211.27 MB	1
5	RUN /bin/sh -c groupadd --gid 1000 node && useradd --uid 1000 --gid node --shell /bin/bash --cre...	3.32 KB	1
6	ENV NODE_VERSION=22.9.0	0 B	1
7	RUN /bin/sh -c ARCH= && dpkgArch="\$(dpkg --print-architecture)" && case "\$dpkgArch#" in i...	53.86 MB	1
8	ENV YARN_VERSION=1.22.22	0 B	1
9	RUN /bin/sh -c set -ex && export GNUPGHOME="\$(mktemp -d)" && for key in 6A010C516600659...	1.25 MB	1
10	COPY docker-entrypoint.sh /usr/local/bin/ # buildkit	444 B	1

(Source: DockerHub)



4. OPTIMIZE IMAGES

Minimizing the Number of Layers:

- Combine multiple RUN commands into one (e.g., chaining commands using &&).
- Removing unnecessary files during the build process (e.g., apt-get clean).
- Avoid using COPY to add files that won't be needed during runtime.

Using Smaller Base Images:

- Alpine Linux vs. Ubuntu: Trade-offs between size and functionality.
- Multi-stage builds: Using one image for building the application and another for the final image (stripped of development tools).

Caching Layers Efficiently:

- Take advantage of layer caching by ordering Dockerfile instructions logically.
- Place less frequently changed commands earlier in the Dockerfile to maximize caching.

Removing Unnecessary Packages and Files:

- Remove temporary files, build dependencies, or logs to reduce image size.
- Clean up the file system after installations (e.g., `rm -rf /var/lib/apt/lists/*`).



4. OPTIMIZE IMAGES - DISTROLESS

- If every change in the file system creates a new layer. Can there be an empty image?
- Yes! Introducing **FROM SCRATCH**
- Images built using this base are called **distroless**.
- By reducing the image size, you are also reducing the **attack surface** for possible vulnerabilities.



4. OPTIMIZE IMAGES - DISTROLESS

- Even though there is a considerable increase in security and performance, images built this way introduce a considerable **development overhead**.
- For learning purposes they are pretty useful, but for production ready environment maybe not the best.
- Google realized this some years back and developed their own **suite of distroless images** (link to Google Distroless repository).

```
# use ubuntu as base image
FROM ubuntu as build-env

# install build-essential package to compile the source code
RUN apt update && apt install -y build-essential

# copy the source code
COPY hello.c hello.c

# Compile and generate binary
RUN gcc -o helloWorld hello.c

# use distroless container to run the program
FROM SCRATCH

# copy the required libraries to run the program
COPY --from=build-env /lib/aarch64-linux-gnu/libc.so.6 /lib/aarch64-linux-gnu/libc.so.6
COPY --from=build-env /lib/ld-linux-aarch64.so.1 /lib/ld-linux-aarch64.so.1

# copy the /etc/passwd
COPY --from=build-env /etc/passwd /etc/passwd

# tmp directory
COPY --from=build-env /tmp /tmp

# copy binary executable to new layer
COPY --from=build-env ./helloWorld ./helloWorld
```

(Source: own)



4. OPTIMIZE IMAGES - DISTROLESS

- In the quest for optimizing images, some projects have appeared.
- Going back to the introduction and understanding the need of having dedicated time for learning for companies that want to use Cloud Native technologies.
- Ubuntu Chiselled images ([link](#)).
- COPA project ([link](#)).



5. CONCLUSSION

Key Takeaways

- Importance of understanding image layers and their role in optimization.
- Best practices for reducing image size and improving efficiency.
- Tools and techniques for analyzing and improving images.

Q&A ?