





# PROTEGER LAS CARGAS DE TRABAJO DE KUBERNETES: DEL CÓDIGO AL CLÚSTER

XOPS CONFERENCE 2024

ÁLVARO REVUELTA



# ABOUT ME

Alvaro Revuelta.

- Systems Developer
- Laboratory for Life Sciences - SciLifeLab






# PRESENTATION STRUCTURE

1. Introduction
2. Inner Loop
3. Outer Loop
4. Other Security challenges
5. Other tools for security
6. Final



# INTRODUCTION - ASSUMPTIONS


- Talk will cover both, developer and platform. But focus a bit more in security from a developer perspective.
- Assumes knowledge about containers and DevOps principles. Some knowledge about Kubernetes is also useful.





# INTRODUCTION - WHY KUBERNETES

- Container management manually is difficult, leading to the rise of orchestration platforms like Kubernetes.
- It is a highly flexible tool, thus, it doesn't make any assumptions and the engineers have to be aware of all the possible vulnerabilities.






# INTRODUCTION - WHY KUBERNETES


Perception on security issues in Kubernetes are underestimated.

- **53 %** of Organizations using Kubernetes experienced security issues; **55%** delayed deployments  
*[redhat.com/rhdc/managed-files/cl-state-of-kubernetes-security-report-2022-ebook-f31209-202205-en.pdf](https://redhat.com/rhdc/managed-files/cl-state-of-kubernetes-security-report-2022-ebook-f31209-202205-en.pdf).*
- **63%** of public code templates had improper configurations; **96%** of cloud applications had known vulnerabilities,  
*[unit42.paloaltonetworks.com/cloud-threat-report-2h-2021](https://unit42.paloaltonetworks.com/cloud-threat-report-2h-2021).*





# INTRODUCTION - INNER OUTER DEVELOPMENT LOOP



SOURCE: REDHAT






## INNER LOOP

1. Write Code.
  2. Build (containerize the application).
  3. Debug Locally.
- This is where the Developer spends most of their time.



## INNER LOOP: KEY CONSIDERATIONS

- Use a secure base image.
- Scan code dependencies for vulnerabilities (e.g., Trivy, Snyk).
- Optimize the application builds by doing good use of concepts such as the cache of layers.
- Increase productivity with Editor/IDE integration: VSCode Remote Containers



(Source: UnionFS : A File System of a Container)



## INNER LOOP: SCAN IMAGE

*trivy image imagetag:version*

```
/ # trivy image rv0lt/flaskrediswebapp:basic | head
2023-04-21T18:04:09.297Z      INFO    Vulnerability scanning is en
2023-04-21T18:04:09.297Z      INFO    Secret scanning is enabled
2023-04-21T18:04:09.297Z      INFO    If your scanning is slow, pl
2023-04-21T18:04:09.297Z      INFO    Please see also https://aqua
ret detection
2023-04-21T18:04:11.456Z      INFO    Detected OS: debian
2023-04-21T18:04:11.456Z      INFO    Detecting Debian vulnerabili
2023-04-21T18:04:11.691Z      INFO    Number of language-specific
2023-04-21T18:04:11.691Z      INFO    Detecting python-pkg vulnera

rv0lt/flaskrediswebapp:basic (debian 10.13)
=====
Total: 146 (UNKNOWN: 5, LOW: 86, MEDIUM: 21, HIGH: 32, CRITICAL: 2)
```



([github.com/aquasecurity/trivy](https://github.com/aquasecurity/trivy))



# INNER LOOP: OPTIMIZE IMAGES

## Minimizing the Number of Layers:

- Combine multiple RUN commands into one (e.g., chaining commands using &&).
- Removing unnecessary files during the build process (e.g., apt-get clean).
- Avoid using COPY to add files that won't be needed during runtime.

## Using Smaller Base Images:

- Alpine Linux vs. Ubuntu: Trade-offs between size and functionality.
- Multi-stage builds: Using one image for building the application and another for the final image (stripped of development tools).

## Caching Layers Efficiently:

- Take advantage of layer caching by ordering Dockerfile instructions logically.
- Place less frequently changed commands earlier in the Dockerfile to maximize caching.

## Removing Unnecessary Packages and Files:

- Remove temporary files, build dependencies, or logs to reduce image size.
- Clean up the file system after installations (e.g., `rm -rf /var/lib/apt/lists/*`).



# INNER LOOP: OPTIMIZE IMAGES

Subsequent building times: +60s

```
Dockerfile basic X Dockerfile basic2 X Do
basic > Dockerfile > ...
1 # use ubuntu as base image
2 FROM ubuntu
3
4 # copy the source code
5 COPY hello.c hello.c
6
7 # install build-essential package to compi
8 RUN apt update
9 RUN apt install -y build-essential
10
11 # Compile and generate binary
12 RUN gcc -o helloWorld hello.c
13
14 # Run the program
15 ENTRYPOINT ["/helloWorld"]
16
17
```

Subsequent building times: 1s

```
Dockerfile basic Dockerfile basic2 X Do
basic2 > Dockerfile > ...
1 # use ubuntu as base image
2 FROM ubuntu
3
4 # install build-essential package to compi
5 RUN apt update
6 RUN apt install -y build-essential
7
8 # copy the source code
9 COPY hello.c hello.c
10
11 # Compile and generate binary
12 RUN gcc -o helloWorld hello.c
13
14 # Run the program
15 ENTRYPOINT ["/helloWorld"]
16
17
```



# INNER LOOP: OPTIMIZE IMAGES

```
# use ubuntu as base image
FROM ubuntu as build-env

# install build-essential package to compile the source code
RUN apt update && apt install -y build-essential

# copy the source code
COPY hello.c hello.c

# Compile and generate binary
RUN gcc -o helloWorld hello.c

# FROM alpine for an even greater size reduce
FROM ubuntu

# copy binary executable to new layer
COPY --from=build-env ./helloWorld ./helloWorld

# Run the program
ENTRYPOINT ["/helloWorld"]
```

WEIGHT IN MB REDUCED BY 4X TIMES



## OUTER LOOP

1. Code Review
  2. Automated testing
  3. Production build
  4. Compliance, and security checks,
  5. Deployment to target environments.
- The platform engineer sets up the automation processes.



## OUTER LOOP: CI - CD

**Continuous Integration:**  
Build-Test-Merge.

**Continuous Delivery:**  
Someone decides when to  
push to prod.

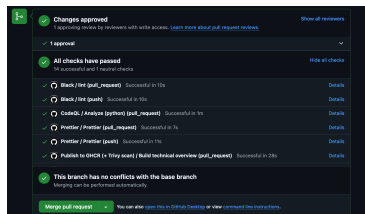
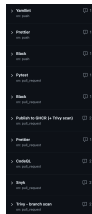
**Continuous Deployment:**  
Automatically deploy to prod.






## OUTER LOOP: SECURITY IN CI/CD

- Scan code and images during builds.
- Block deployments for high-severity vulnerabilities.





## ArgoCD





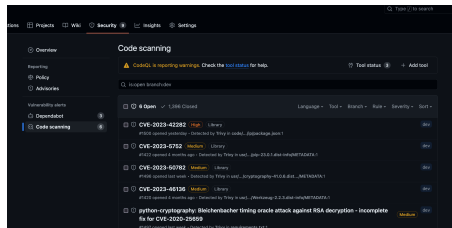


# OUTER LOOP: CONTINUOUS VULNERABILITY MANAGEMENT

## Why Continuous Scanning Matters

- The evolving threat landscape.
- New Vulnerabilities pop up quite often, CVE database

Integrate tools like Trivy into CI/CD pipelines to schedule scans of image repositories.





## OTHER SECURITY CHALLENGES IN K8S

### Default permissions

- By default, containers run as a root user.
- Malicious agents can exploit root access.
- We need to define security contexts to run as a non-privileged user and apply restrictions.

### Communication between pods

- By default, all resources can communicate with each other.
- We need to ensure resource separation to provide better security in case of breach.
- The solution is to limit traffic flow for better control.



# VERY INTERESTING LECTURE



National Security Agency  
Cybersecurity and Infrastructure Security Agency

Cybersecurity Technical Report

## **Kubernetes Hardening Guide**

KUBERNETS HARDENING GUIDE



## OTHER TOOLS FOR SECURITY

- Kube-hunter: Cluster Misconfiguration Detection; Pen-testing.
  - Same makers as Trivy
  - [github.com/aquasecurity/kube-hunter](https://github.com/aquasecurity/kube-hunter)
- Kubescape: similar than kube-hunter.
  - [github.com/kubescape](https://github.com/kubescape)
- Kube-linter: Same as Trivy, but to find misconfiguration in k8s yaml files.
- Falco: Runtime Security.

Vulnerabilities  
For further information about a vulnerability, search its ID in:  
<https://avd.aquasec.com/>

ID	LOCATION	MITRE CATEGORY	VULNERABILITY	DESCRIPTION	EVIDENCE
None	Local to Pod (kube-hunter-075v6)	Lateral Movement // ARP poisoning and IP spoofing	CAP_NET_RAW Enabled	CAP_NET_RAW is enabled by default for pods. If an attacker manages to compromise a pod, they could potentially take advantage of this capability to perform network attacks on other pods running on the same node	
KHWR02	10.96.0.1:443	Initial Access // Exposed sensitive interfaces	K8s Version Disclosure	The kubernetes version could be obtained from the /version endpoint	v1.27.2





# CNCF LANDSCAPE



OFFICIAL COMPLIANT KUBERNETES CNCF.IO



# FINAL

CONTACT:

LINKEDIN - ALVARO REVUELTA MARTINEZ  
ALVARO.REVUELTA@SCILIFELAB.UU.SE



# FINAL

Q & A